

Sahil Y. Tike
BE B3 (212066)

Transformers From Scratch Using Pytorch

Type *Markdown* and LaTeX: α^2

1. Introduction

Implement transformers in "Attention is all you need paper" from scratch using Pytorch. Basically transformer have an encoder-decoder architecture. It is common for language translation models.

2. Import Libraries

```
] pip install torchvision
```

```
In [ ]: Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied:
3.8/site-packages (0.15.1)
Collecting torch==2.0.0
Note: you may need to restart the kernel to use updated packages.
```

```
] pip install torchtext==0.10.0
```

```
In [ ]: Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: torchtext==0.10.0 in /home/dipali/.local/lib/python3.8/site-packages (0.10.0)
Requirement already satisfied: torch==1.9.0 in /home/dipali/.local/lib/python3.8/site-packages (from torchtext==0.10.0) (1.9.0)
Requirement already satisfied: numpy in /home/dipali/.local/lib/python3.8/site-packages (from torchtext==0.10.0) (1.24.1)
Requirement already satisfied: requests in /usr/lib/python3/dist-packages (from torchtext==0.10.0) (2.22.0)
Requirement already satisfied: tqdm in /home/dipali/.local/lib/python3.8/site-packages (from torchtext==0.10.0) (4.64.1)
Requirement already satisfied: typing-extensions in /home/dipali/.local/lib/python3.8/site-packages (from torch==1.9.0->torchtext==0.10.0) (4.4.0)

[notice] A new release of pip is available: 23.0 -> 23.0.1
[notice] To update, run: python3 -m pip
Note: you may need to restart the kernel to
```

new

```
install --upgrade pip
use updated packages.
```

```

]
In [ ]: # importing required libraries
import torch.nn as nn
import torch
import torch.nn.functional as F
import math, copy, re
import warnings
import pandas as pd
import numpy as np
import seaborn as sns
import torchtext
import matplotlib.pyplot as plt
warnings.simplefilter("ignore")
print(torch.__version__)

```

1.9.0+cu102

3. Basic components

```

In [ ]: class Embedding(nn.Module):
    def __init__(self, vocab_size, embed_dim):
        """
        Args:
            vocab_size: size of vocabulary
            embed_dim: dimension of embeddings
        """
        super(Embedding, self).__init__()
        self.embed = nn.Embedding(vocab_size, embed_dim)
    def forward(self, x):
        """
        Args:
            x: input vector
        Returns:
            out: embedding vector
        """
        out = self.embed(x)
        return out

```

```

]
In [ : # register buffer in Pytorch ->
# If you have parameters in your model, which should be saved and restored in
# but not trained by the optimizer, you should register them as buffers.

class PositionalEmbedding(nn.Module):
    def __init__(self,max_seq_len,embed_model_dim):
        """
        Args:
            seq_len: length of input sequence
            embed_model_dim: demension of embedding
        """
        super(PositionalEmbedding, self).__init__()
        self.embed_dim = embed_model_dim

        pe = torch.zeros(max_seq_len,self.embed_dim)
        for pos in range(max_seq_len):
            for i in range(0,self.embed_dim,2):
                pe[pos, i] = math.sin(pos / (10000 ** ((2 * i)/self.embed_dim)))
                pe[pos, i + 1] = math.cos(pos / (10000 ** ((2 * (i + 1))/self.
pe = pe.unsqueeze(0)
self.register_buffer('pe', pe)

    def forward(self, x):
        """
        Args:
            x: input vector
        Returns:
            x: output
        """

        # make embeddings relatively larger
        x = x * math.sqrt(self.embed_dim)
        #add constant to embedding
        seq_len = x.size(1)
        x = x + torch.autograd.Variable(self.pe[:, :seq_len], requires_grad=False)
        return x

```

```

In [ ]: class MultiHeadAttention(nn.Module):
    def __init__(self, embed_dim=512, n_heads=8):
        """
        Args:
            embed_dim: dimension of embedding vector output
            n_heads: number of self attention heads
        """
        super(MultiHeadAttention, self).__init__()

        self.embed_dim = embed_dim      #512 dim
        self.n_heads = n_heads          #8
        self.single_head_dim = int(self.embed_dim / self.n_heads)      #512/8 = 64

        #key, query and value matrixes      #64 x 64
        self.query_matrix = nn.Linear(self.single_head_dim, self.single_head_dim)
        self.key_matrix = nn.Linear(self.single_head_dim, self.single_head_dim)
        self.value_matrix = nn.Linear(self.single_head_dim, self.single_head_dim)
        self.out = nn.Linear(self.n_heads*self.single_head_dim, self.embed_dim)

    def forward(self, key, query, value, mask=None):      #batch_size x sequence_length
        """
        Args:
            key : key vector
            query : query vector
            value : value vector
            mask: mask for decoder

        Returns:
            output vector from multihead attention
        """
        batch_size = key.size(0)
        seq_length = key.size(1)

        # query dimension can change in decoder during inference.
        # so we cant take general seq_length
        seq_length_query = query.size(1)

        # 32x10x512
        key = key.view(batch_size, seq_length, self.n_heads, self.single_head_dim)
        query = query.view(batch_size, seq_length_query, self.n_heads, self.single_head_dim)
        value = value.view(batch_size, seq_length, self.n_heads, self.single_head_dim)

        k = self.key_matrix(key)      # (32x10x8x64)
        q = self.query_matrix(query)
        v = self.value_matrix(value)

        q = q.transpose(1,2) # (batch_size, n_heads, seq_len, single_head_dim)
        k = k.transpose(1,2) # (batch_size, n_heads, seq_len, single_head_dim)
        v = v.transpose(1,2) # (batch_size, n_heads, seq_len, single_head_dim)

        # computes attention
        # adjust key for matrix multiplication
        k_adjusted = k.transpose(-1,-2) #(batch_size, n_heads, single_head_dim, seq_len)
        product = torch.matmul(q, k_adjusted) #(32 x 8 x 10 x 64) x (32 x 8 x 10 x 64)

```

```

# fill those positions of product matrix as (-1e20) where mask position
if mask is not None:
    product = product.masked_fill(mask == 0, float("-1e20"))

#dividing by square root of key dimension
product = product / math.sqrt(self.single_head_dim) # / sqrt(64)

#applying softmax
scores = F.softmax(product, dim=-1)

#mutiply with value matrix
scores = torch.matmul(scores, v) ##(32x8x 10x 10) x (32 x 8 x 10 x 64)

#concatenated output
concat = scores.transpose(1,2).contiguous().view(batch_size, seq_length, self.single_head_dim)

output = self.out(concat) #(32,10,512) -> (32,10,512)

return output

```

4. Encoder

```

In [ ]: class TransformerBlock(nn.Module):
        def __init__(self, embed_dim, expansion_factor=4, n_heads=8):
            super(TransformerBlock, self).__init__()

            """
            Args:
                embed_dim: dimension of the embedding
                expansion_factor: factor which determines output dimension of linear
                n_heads: number of attention heads

            """
            self.attention = MultiHeadAttention(embed_dim, n_heads)

            self.norm1 = nn.LayerNorm(embed_dim)
            self.norm2 = nn.LayerNorm(embed_dim)

            self.feed_forward = nn.Sequential(
                nn.Linear(embed_dim, expansion_factor*embed_dim),
                nn.ReLU(),
                nn.Linear(expansion_factor*embed_dim, embed_dim)
            )

            self.dropout1 = nn.Dropout(0.2)
            self.dropout2 = nn.Dropout(0.2)

        def forward(self, key, query, value):
            """
            Args:
                key: key vector
                query: query vector
                value: value vector
                norm2_out: output of transformer block

            """

            attention_out = self.attention(key, query, value) #32x10x512
            attention_residual_out = attention_out + value #32x10x512
            norm1_out = self.dropout1(self.norm1(attention_residual_out)) #32x10x512

            feed_fwd_out = self.feed_forward(norm1_out) #32x10x512 -> #32x10x2048
            feed_fwd_residual_out = feed_fwd_out + norm1_out #32x10x512
            norm2_out = self.dropout2(self.norm2(feed_fwd_residual_out)) #32x10x512

            return norm2_out

class TransformerEncoder(nn.Module):
    """
    Args:
        seq_len : length of input sequence
        embed_dim: dimension of embedding
        num_layers: number of encoder layers
        expansion_factor: factor which determines number of linear layers in f
        n_heads: number of heads in multihead attention
    """

```

```

Returns:
    out: output of the encoder
"""
def __init__(self, seq_len, vocab_size, embed_dim, num_layers=2, expansion_factor=2):
    super(TransformerEncoder, self).__init__()

    self.embedding_layer = Embedding(vocab_size, embed_dim)
    self.positional_encoder = PositionalEncoding(seq_len, embed_dim)

    self.layers = nn.ModuleList([TransformerBlock(embed_dim, expansion_factor)
                                   for _ in range(num_layers)])

def forward(self, x):
    embed_out = self.embedding_layer(x)
    out = self.positional_encoder(embed_out)
    for layer in self.layers:
        out = layer(out, out, out)

    return out  #32x10x512

```


5. Decoder

```

In [ ]: class DecoderBlock(nn.Module):
    def __init__(self, embed_dim, expansion_factor=4, n_heads=8):
        super(DecoderBlock, self).__init__()

        """
        Args:
            embed_dim: dimension of the embedding
            expansion_factor: factor which determines output dimension of linear
            n_heads: number of attention heads

        """
        self.attention = MultiHeadAttention(embed_dim, n_heads=8)
        self.norm = nn.LayerNorm(embed_dim)
        self.dropout = nn.Dropout(0.2)
        self.transformer_block = TransformerBlock(embed_dim, expansion_factor,

    def forward(self, key, query, x, mask):

        """
        Args:
            key: key vector
            query: query vector
            value: value vector
            mask: mask to be given for multi head attention
        Returns:
            out: output of transformer block

        """

        #we need to pass mask mask only to fst attention
        attention = self.attention(x,x,x,mask=mask) #32x10x512
        value = self.dropout(self.norm(attention + x))

        out = self.transformer_block(key, query, value)

        return out

class TransformerDecoder(nn.Module):
    def __init__(self, target_vocab_size, embed_dim, seq_len, num_layers=2, ex
        super(TransformerDecoder, self).__init__()

        """
        Args:
            target_vocab_size: vocabulary size of target
            embed_dim: dimension of embedding
            seq_len : length of input sequence
            num_layers: number of encoder layers
            expansion_factor: factor which determines number of linear layers i
            n_heads: number of heads in multihead attention

        """

        self.word_embedding = nn.Embedding(target_vocab_size, embed_dim) 10
        self.position_embedding = PositionalEmbedding(seq_len, embed_dim)

        self.layers = nn.ModuleList(

```

```

        [
            DecoderBlock(embed_dim, expansion_factor=4, n_heads=8)
            for _ in range(num_layers)
        ]

    )
    self.fc_out = nn.Linear(embed_dim, target_vocab_size)
    self.dropout = nn.Dropout(0.2)

    def forward(self, x, enc_out, mask):
        """
        Args:
            x: input vector from target
            enc_out : output from encoder layer
            trg_mask: mask for decoder self attention
        Returns:
            out: output vector
        """

        x = self.word_embedding(x) #32x10x512
        x = self.position_embedding(x) #32x10x512
        x = self.dropout(x)

        for layer in self.layers:
            x = layer(enc_out, x, enc_out, mask)

        out = F.softmax(self.fc_out(x))

        return out

```

Finally we will arrange all submodules and creates the entire tranformer architecture.

In []:

```
class Transformer(nn.Module):
    def __init__(self, embed_dim, src_vocab_size, target_vocab_size, seq_length):
        super(Transformer, self).__init__()

        """
        Args:
            embed_dim: dimension of embedding
            src_vocab_size: vocabulary size of source
            target_vocab_size: vocabulary size of target
            seq_length : length of input sequence
            num_layers: number of encoder layers
            expansion_factor: factor which determines number of linear layers in each block
            n_heads: number of heads in multihead attention
        """

        self.target_vocab_size = target_vocab_size

        self.encoder = TransformerEncoder(seq_length, src_vocab_size, embed_dim)
        self.decoder = TransformerDecoder(target_vocab_size, embed_dim, seq_length)

    def make_trg_mask(self, trg):
        """
        Args:
            trg: target sequence
        Returns:
            trg_mask: target mask
        """
        batch_size, trg_len = trg.shape
        # returns the lower triangular part of matrix filled with ones
        trg_mask = torch.tril(torch.ones((trg_len, trg_len))).expand(
            batch_size, 1, trg_len, trg_len
        )
        return trg_mask

    def decode(self, src, trg):
        """
        for inference
        Args:
            src: input to encoder
            trg: input to decoder
        out:
            out_labels : returns final prediction of sequence
        """
        trg_mask = self.make_trg_mask(trg)
        enc_out = self.encoder(src)
        out_labels = []
        batch_size, seq_len = src.shape[0], src.shape[1]
        #outputs = torch.zeros(seq_len, batch_size, self.target_vocab_size)
        out = trg
        for i in range(seq_len): #10
            out = self.decoder(out, enc_out, trg_mask) #bs x seq_len x vocab_dim
            # taking the last token
            out = out[:, -1, :]

```

```

        out = out.argmax(-1)
        out_labels.append(out.item())
        out = torch.unsqueeze(out,axis=0)

    return out_labels

def forward(self, src, trg):
    """
    Args:
        src: input to encoder
        trg: input to decoder
    out:
        out: final vector which returns probabilities of each target word
    """
    trg_mask = self.make_trg_mask(trg)
    enc_out = self.encoder(src)

    outputs = self.decoder(trg, enc_out, trg_mask)
    return outputs

```

6. Testing Code

Suppose we have input sequence of length 10 and target sequence of length 10.

```

In [ ]: src_vocab_size = 11
        target_vocab_size = 11
        num_layers = 6
        seq_length= 12

        # let 0 be sos token and 1 be eos token
        src = torch.tensor([[0, 2, 5, 6, 4, 3, 9, 5, 2, 9, 10, 1],
                             [0, 2, 8, 7, 3, 4, 5, 6, 7, 2, 10, 1]])
        target = torch.tensor([[0, 1, 7, 4, 3, 5, 9, 2, 8, 10, 9, 1],
                                [0, 1, 5, 6, 2, 4, 7, 6, 2, 8, 10, 1]])

        print(src.shape,target.shape)
        model = Transformer(embed_dim=512, src_vocab_size=src_vocab_size,
                             target_vocab_size=target_vocab_size, seq_length=seq_length,
                             num_layers=num_layers, expansion_factor=4, n_heads=8)
        model

        torch.Size([2, 12]) torch.Size([2, 12])

```

```

Out[10]: Transformer(
  (encoder): TransformerEncoder(
    (embedding_layer): Embedding(
      (embed): Embedding(11, 512)
    )
    (positional_encoder): PositionalEmbedding()
    (layers): ModuleList(
      (0): TransformerBlock(
        (attention): MultiHeadAttention(
          (query_matrix): Linear(in_features=64, out_features=64, bias=False)
          (key_matrix): Linear(in_features=64, out_features=64, bias=False)
          (value_matrix): Linear(in_features=64, out_features=64, bias=False)
          (out): Linear(in_features=512, out_features=512, bias=True)
        )
        (norm1): LayerNorm((512), eps=1e-05, elementwise_affine=True)

```

```

In [ ]: out = model(src, target)
        out.shape

```

```

Out[11]: torch.Size([2, 12, 11])

```

In []:

```
# inference
model = Transformer(embed_dim=512, src_vocab_size=src_vocab_size,
                    target_vocab_size=target_vocab_size, seq_length=seq_length,
                    num_layers=num_layers, expansion_factor=4, n_heads=8)

src = torch.tensor([[0, 2, 5, 6, 4, 3, 9, 5, 2, 9, 10, 1]])
trg = torch.tensor([[0]])
print(src.shape, trg.shape)
out = model.decode(src, trg)
out

torch.Size([1, 12]) torch.Size([1, 1])
```

Out[12]: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]